

Language Modeling with Torch

Nicholas Leonard

Element Inc.

<https://github.com/nicholas-leonard>



torch

Agenda

- ▶ Meeting Torch
 - ▶ What is Torch?
 - ▶ When should I use Torch?
 - ▶ How does Torch work?
- ▶ Language Modeling with Torch
 - ▶ Understanding Language Models
 - ▶ Recurrent Neural Networks
 - ▶ Generating sentences



Meeting Torch

Lua, tensors and neural networks



torch

What is Torch?

Lua, not Python

► Examples of Lua types :

String

```
a = 'hello'  
string.sub(a, 'he[l]+)', 'jell')  
print(a)  
jello
```

Number

```
b = 3  
print(b + 4)  
7
```

Table

```
c = {1,2,3,key='value'}  
c.key = 4  
c[2] = 7  
print(c)  
{1,7,3,key=4}
```

- Lua is lightweight, elegant and easy to learn
- Easy to interface with C/CUDA
- Fast for-loops
- Functions as first class citizens, closures



torch

What is Torch?

Tensors for linear algebra

$$\begin{matrix} \begin{pmatrix} v_1 \\ v_2 \\ v_3 \\ \vdots \\ v_n \end{pmatrix} & \begin{pmatrix} v_{11}, v_{12}, v_{13}, \dots, v_{1m} \\ v_{21}, v_{22}, v_{23}, \dots, v_{2m} \\ v_{31}, v_{32}, v_{33}, \dots, v_{3m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ v_{n1}, v_{n2}, v_{n3}, \dots, v_{nm} \end{pmatrix} \\ \text{Vector} & \text{Matrix} \end{matrix}$$

- ▶ Tensors are N-dimension arrays

- ▶ Initialize and matrix-matrix multiply two tensors:

<code>a = torch.Tensor(2,3)</code>	<code>b = torch.Tensor({1},{2},{3})</code>	<code>c = torch.mm(a,b)</code>
<code>a:uniform(0,1)</code>	<code>print(b)</code>	<code>c:mm(a,b)</code>
<code>print(a)</code>	1	<code>c = a * b</code>
0.6187 0.8752 0.1496	2	
0.6979 0.8230 0.1439	3	
[torch.DoubleTensor of size 2x3]	[torch.DoubleTensor of size 3x1]	

- ▶ Basic linear algebra sub-routines (BLAS)

- ▶ Sub-tensor extraction, etc.

- ▶ Can be run on GPU/CPU

- ▶ Different types : DoubleTensor, FloatTensor, LongTensor, CudaTensor, etc.



torch

What is Torch?

Neural Networks for deep learning

- ▶ Train neural networks using back-propagation (gradient descent + chain rule)

-- model

```
mlp = nn.Sequential()  
:add(nn.Linear(3,4))  
:add(nn.Tanh())  
:add(nn.Linear(4,2))  
:add(nn.Sigmoid())
```

-- loss function

```
mse = nn.MSECriterion()
```

-- data sample

```
input = torch.randn(3)  
target = torch.randn(2)
```

-- forward

```
output = mlp:forward(input)  
loss = mse:forward(output, target)
```

-- backward

```
gradOutput = mse:backward(output, target)  
mlp:zeroGradParameters()  
mlp:backward(input, gradOutput)
```

-- update

```
mlp:updateParameters(0.1)
```

- ▶ Automatic gradient differentiation
 - ▶ No compilation (TensorFlow, Theano)
 - ▶ Imperative programming (easier to debug than symbolic graphs)
- ▶ Extend by implementing nn.Modules or nn.Criterions



When should Torch be used?

Research and production

- ▶ Deep learning research
 - ▶ Assemble existing Modules and Criteria
 - ▶ Or implement your own
 - ▶ Easy to debug (no compilation or symbolic graph)
 - ▶ Lots of packages to pick from (torch, nn, optim, image, dataloader, rnn, etc.)
- ▶ Production environment
 - ▶ Lua was designed for embedded systems
 - ▶ Servers (GPU/CPU, Ubuntu/Linux, Mac OS X)
 - ▶ Smart phones (Android, iOS)
- ▶ Fast execution
 - ▶ glample/rnn-benchmarks
 - ▶ soumith/convnet-benchmarks




torch


How does Torch work?

Mental model

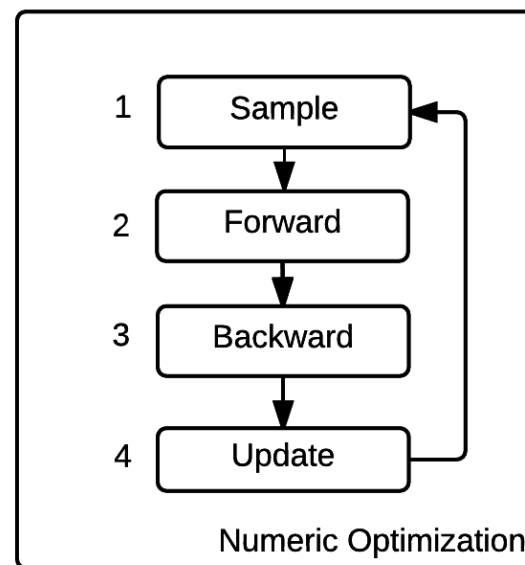
Prepare Dataset



Define Model and Loss



Train Model



torch

Language Models

Generating sentences using recurrent neural networks



torch

Understanding language models

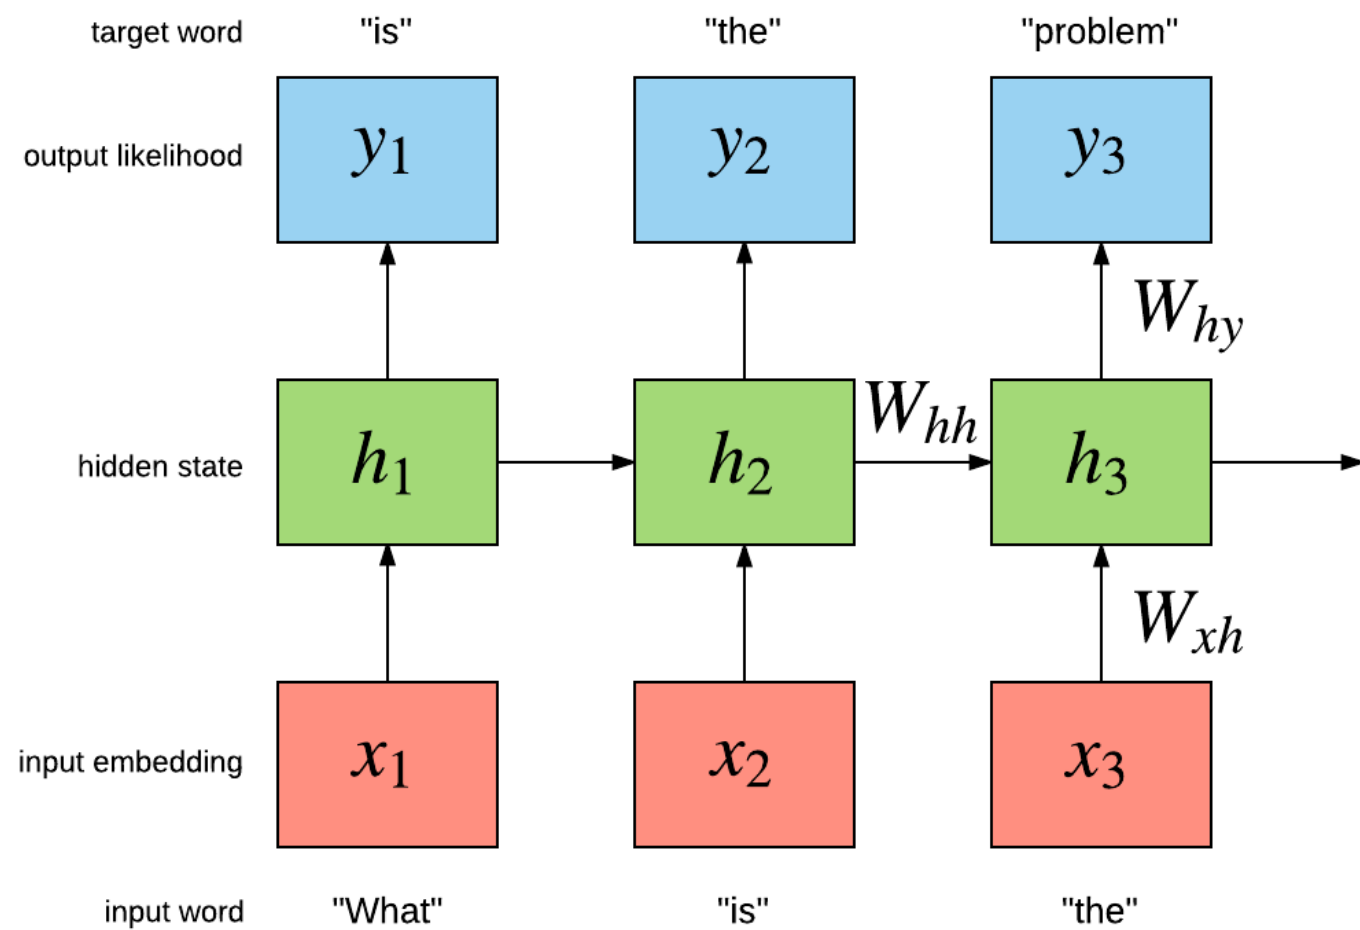
- ▶ Language model dataset is a corpus of text samples
 - ▶ “Why does it always rain on Saturdays?”
- ▶ Maximize the likelihood of a word given previous words: $P(w_t | w_1, \dots, w_{t-1})$
- ▶ Applications:
 - ▶ Generate sentences (sample next word given previous words)
 - ▶ Auto-complete
 - ▶ Measure the probability of a sentence (e.g. does this translation make sense?)

$$P(w_1, w_2, \dots, w_T) = \prod_{t=1}^T P(w_t | w_1, \dots, w_{t-1})$$



torch

Recurrent Neural Networks



torch

Recurrent Neural Networks

- ▶ Model architectures
 - ▶ Simple Recurrent Neural Networks (Simple RNN)
 - ▶ Long Short Term Memory (LSTM)
 - ▶ Gated Recurrent Units (GRU)
- ▶ Used to model sequential data
 - ▶ Video frames, text, time-series
- ▶ Different applications
 - ▶ Language modeling, sentiment analysis, machine translation
- ▶ Element-Research/rnn package provides RNNs for Torch
 - ▶ Tutorials, documentation, models, criterions and examples



torch

Generating sentences

- ▶ **Recursively sample sentences (one word at a time)**

<S> The company said its net profit rose to \$ 289 million , or 96 cents per share , in the three months ended on March 31 compared with \$ 173 million , or \$ 0.68 a share , a year ago . </S>

<S> But I 've been a bit disappointed with our performance , " said Wenger . </S>

<S> The first is an even bigger problem . </S>

<S> The next big thing for him is he will be able to tell the world he is thinking about his future . </S>

<S> The new rules have been added to the legislation so that they don 't have to be approved for public use . </S>

<S> The Pentagon 's top counter-terrorism official , who has been in charge of a new system of intelligence collection and inspection , wrote in an e-mail message that while the new system could be easily implemented , it remains an option . </S>

<S> " I was trying to get a glass of water . </S>

<S> Later he was driven to a nearby house where he was later found to be severely ill . </S>

- ▶ **40GB model learns to generate coherent sentences**

- ▶ **Trained on Google 1-billion words dataset**

Multi-GPU Multi-Layer LSTM

```
nn.Serial @ nn.Sequential {
  [input -> (1) -> (2) -> (3) -> output]
  (1): nn.ParallelTable {
    input
    \-> (1): nn.Sequential {
      [input -> (1) -> (2) -> (3) -> (4) -> (5) -> (6) -> (7) -> (8) -> (9) -> (10) -> (11) -> (12) -> output]
      (1): nn.Convert
      (2): nn.GPU(2) @ nn.Concat {
        input
        | \-> (1): nn.GPU(1) @ nn.LookupTableMaskZero
        | \-> (2): nn.GPU(2) @ nn.LookupTableMaskZero
        ... -> output
      }
      (3): nn.GPU(2) @ nn.Dropout(0.2, busy)
      (4): nn.GPU(1) @ nn.SeqLSTM
      (5): nn.GPU(1) @ nn.Dropout(0.2, busy)
      (6): nn.GPU(1) @ nn.SeqLSTM
      (7): nn.GPU(1) @ nn.Dropout(0.2, busy)
      (8): nn.GPU(2) @ nn.SeqLSTM
      (9): nn.GPU(2) @ nn.Dropout(0.2, busy)
      (10): nn.GPU(2) @ nn.SeqLSTM
      (11): nn.GPU(2) @ nn.Dropout(0.2, busy)
      (12): nn.GPU(3) @ nn.SplitTable
    }
    \-> (2): nn.Identity
    ... -> output
  }
  (2): nn.ZipTable
  (3): nn.GPU(3) @ nn.Sequencer @ nn.Recursor @ nn.MaskZero @ nn.NCEModule(2048 -> 793471)
}
```



torch

For more info:

<http://torch.ch>

<https://github.com/Element-Research/rnn>



torch